



NaN.builders model routing

Cheat sheet for the **agentic-workflow** skills — which NaN Cloud model to run, primary + fallback, Thinking on/off, and effort level. Sign up: cloud.nan.builders/r/7GK06FX8

Per-skill routing — primary & fallback

Skill	Primary	Think/Effort	Fallback	Think/Effort
init-workspace	GLM-5.2	on / High	Mimo V2.5	on / High
plan-feature (+ internals)	GLM-5.2	on / High	Mimo V2.5	on / High
plan-fix	GLM-5.2	on / High	Mimo V2.5	on / High
execute-phase	Qwen3.6	off / Medium	DeepSeek V4 Flash	off / Medium
execute-phase (subtle logic)	GLM-5.2	on / High	Mimo V2.5	on / High
review-change	GLM-5.2	on / High	Mimo V2.5	on / High
audit-pr	GLM-5.2	on / High	Mimo V2.5 (never lower)	on / High
product-audit	GLM-5.2	on / Max	Mimo V2.5 (never lower)	on / Max
audit-docs	Qwen3.6	off / Medium	DeepSeek V4 Flash	off / Medium
audit-docs (deep audit)	GLM-5.2	on / High	Mimo V2.5	on / High
triage-issue	GLM-5.2	on / High	Mimo V2.5	on / High
log-session	DeepSeek V4 Flash	off / Low	Gemma4	off / Low
ship-roadmap (conductor)	GLM-5.2	on / High	Mimo V2.5	on / High
ship-roadmap (exec subagents)	Qwen3.6	off / Medium	DeepSeek V4 Flash	off / Medium
bump-skill	Qwen3.6	off / Medium	DeepSeek V4 Flash	off / Low

Quick tasks

Task	Model	Config	Why
Small fix (XS/S) — plan it	GLM-5.2 (ideal) / Qwen3.6 (ok for XS)	on / High	Risk/blast-radius judgment lives here; on an obvious typo, Qwen3.6+thinking is enough — review catches the rest.
Small fix — implement it (execute-phase --fix)	Qwen3.6	off / Medium	Mechanical against an already-scoped SPEC.
Commits & PRs only (messages, git add/push, gh pr create)	DeepSeek V4 Flash	off / Minimal–Low	Pure forge mechanics — no judgment being bought. Gemma4 also works.

When to reach for each model

GLM-5.2 — Anything that ends in a **verdict or design decision**: planning, reviewing a diff, the merge gate, triaging an issue. If the output is PASS/FAIL, MERGE-READY, or a SPEC — this is the model.

Mimo V2.5 — Three jobs: (1) fallback for everything GLM-5.2 does; (2) **cross-family reviewer** when Qwen wrote the code (adds independence); (3) **1M context** — the PRIMARY pick (not just fallback) whenever a review/audit needs to swallow a repo/diff too big for GLM's 256K, regardless of whether GLM is up.

Qwen3.6 — Writing code against a SPEC, doc checks, bookkeeping. 35B with speculative decoding (~2x throughput) — fast, cheap workhorse. Its verdicts gate nothing. Also has vision — fine for mechanical UI checks.

DeepSeek V4 Flash — Cheap + huge context (1M): session logs, grep-shaped evidence gathering, commit/PR mechanics, summarizing a giant log. Never for verdicts.

Gemma4 — Minimum-viable tier: trivial tasks, last-resort small-tier fallback. Also has vision. No published member quota — the safe pick to burn on cheap/bulk work instead of Mimo/Flash.

The rule that ties it together: the model that judges must be $\geq$ the model that wrote. The two merge-gating verdicts — **audit-pr** and **product-audit** — never drop below the GLM-5.2 / Mimo V2.5 pair. If both are down: defer — gate the merge by hand, postpone the product audit. A shallow verdict is worse than a delayed one.

Easy-to-miss cases

Case	What to do
Huge diff/repo (> GLM's 256K context)	Route straight to Mimo V2.5 as the PRIMARY for that pass (1M ctx) — don't wait for GLM to fail first.
UI screenshot check in review-design/review-a11y	Qwen3.6 and Gemma4 have vision too, not just Mimo — use Qwen3.6 if the check is mechanical; reserve Mimo/GLM for when the design judgment itself is what's being reviewed.
Mimo V2.5 / DeepSeek V4 Flash quota (500M tok/member/mo — the only two models with a published cap)	Don't burn Mimo's quota on cheap bulk work — it's the fallback for the two never-degrade verdicts. Prefer Gemma4 (no published quota) for cheap/bulk tasks instead.
Entire NaN cluster down (not just one model)	No NaN fallback exists at that point — fall back to a different provider entirely, or defer merge-gating work until service is back.